

Q1. Assume the input matrix I and weight matrix W. Write a Python code to convolve the weight matrix on image I and apply average pooling without using the inbuilt function. (10)

Input image (I)

1	2	2
1	3	4
1	0	0

Weight matrix(W)

1	-1
1	-1
0	1

Importing libraries and defining input and weight matrices

```
In [7]: import numpy as np
```

```
# Defining input image matrix (I) and weight matrix (W)
I = np.array([[1, 2, 2],
              [1, 3, 4],
              [1, 0, 0]])

W = np.array([[1, -1],
              [1, -1],
              [0, 1]])
```

```
In [8]: # Perform convolution operation
result_convolution = np.zeros((2, 2))
for i in range(2):
    for j in range(2):
        result_convolution[i, j] = np.sum(I[i:i+3, j:j+2] * W[i:i+3, j:j+2])

print("Result after convolution: ", result_convolution)

Result after convolution:  [[ -3. -11.]
 [ -2.  -7.]]
```

```
In [9]: # Perform average pooling
result_pooling = np.mean(result_convolution)

print("Result after average pooling:", result_pooling)

Result after average pooling: -5.75
```

Q2. Take the MNIST dataset and create a CNN architecture to create a classification model. Use Adam optimiser from the library and without a library and comment on your observations. (10)

Implementing CNN architecture without using any deep learning library and use basic numerical computation to create the model.

```

In [1]: import numpy as np
        from tqdm import tqdm
        from scipy.special import logsumexp
        from keras.datasets.mnist import load_data

        class MLP():

            def __init__(self, din, dout):
                self.W = (2 * np.random.rand(dout, din) - 1) * (np.sqrt(6) / np.sqrt(din))
                self.b = (2 * np.random.rand(dout) - 1) * (np.sqrt(6) / np.sqrt(din))

            def forward(self, x): # x.shape = (batch_size, din)
                self.x = x # Storing x for latter (backward pass)
                return x @ self.W.T + self.b

            def backward(self, gradout):
                self.deltaW = gradout.T @ self.x
                self.deltab = gradout.sum(0)
                return gradout @ self.W

        class SequentialNN():

            def __init__(self, blocks: list):
                self.blocks = blocks

            def forward(self, x):

                for block in self.blocks:
                    x = block.forward(x)

                return x

            def backward(self, gradout):

                for block in self.blocks[::-1]:
                    gradout = block.backward(gradout)

                return gradout

        class ReLU():

            def forward(self, x):
                self.x = x
                return np.maximum(0, x)

            def backward(self, gradout):
                new_grad = gradout.copy()
                new_grad[self.x < 0] = 0.
                return new_grad

        class LogSoftmax():

            def forward(self, x):
                self.x = x
                return x - logsumexp(x, axis=1)[..., None]

            def backward(self, gradout):
                gradients = np.eye(self.x.shape[1])[None, ...]
                gradients = gradients - (np.exp(self.x) / np.sum(np.exp(self.x), axis=1)[..., None])
                return (np.matmul(gradients, gradout)[..., None])[..., :, 0]

```

```

class NLLLoss():

    def forward(self, pred, true):
        self.pred = pred
        self.true = true

        loss = 0
        for b in range(pred.shape[0]):
            loss -= pred[b, true[b]]
        return loss

    def backward(self):
        din = self.pred.shape[1]
        jacobian = np.zeros((self.pred.shape[0], din))
        for b in range(self.pred.shape[0]):
            jacobian[b, self.true[b]] = -1

        return jacobian # batch_size x din

    def __call__(self, pred, true):
        return self.forward(pred, true)

class Optimizer():

    def __init__(self, lr, compound_nn: SequentialNN):
        self.lr = lr
        self.compound_nn = compound_nn

    def step(self):

        for block in self.compound_nn.blocks:
            if block.__class__ == MLP:
                block.W = block.W - self.lr * block.deltaw
                block.b = block.b - self.lr * block.deltab

def train(model, optimizer, trainX, trainy, loss_fct = NLLLoss(), nb_epochs=
training_loss = []
for epoch in tqdm(range(nb_epochs)):

    # Sample batch size
    batch_idx = [np.random.randint(0, trainX.shape[0]) for _ in range(ba
    x = trainX[batch_idx]
    target = trainy[batch_idx]

    prediction = model.forward(x) # Forward pass
    loss_value = loss_fct(prediction, target) # Compute the Loss
    training_loss.append(loss_value) # Log Loss
    gradout = loss_fct.backward()
    model.backward(gradout) # Backward pass

    # Update the weights
    optimizer.step()
return training_loss

if __name__ == "__main__":
    # Load and process data
    (trainX, trainy), (testX, testy) = load_data()
    trainX = trainX / 255
    testX = testX / 255
    trainX = trainX.reshape(trainX.shape[0], 28 * 28)

```


Implementing CNN using Deep learning libraries

```
In [1]: import tensorflow as tf
from tensorflow.keras import layers, models
from tensorflow.keras.datasets import mnist

# Load the MNIST dataset
(train_images, train_labels), (test_images, test_labels) = mnist.load_data()
train_images = train_images.reshape((60000, 28, 28, 1)).astype('float32') / 255
test_images = test_images.reshape((10000, 28, 28, 1)).astype('float32') / 255

# Define the CNN architecture
model = models.Sequential([
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.Flatten(),
    layers.Dense(64, activation='relu'),
    layers.Dense(10, activation='softmax')
])

# Compile the model
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

# Train the model
history = model.fit(train_images, train_labels, epochs=5, batch_size=64, validation_data=(test_images, test_labels))

# Evaluate the model
test_loss, test_acc = model.evaluate(test_images, test_labels)
print(f'Test accuracy: {test_acc}')
```

```
Epoch 1/5
750/750 [=====] - 14s 18ms/step - loss: 0.2105 - accuracy: 0.9361 - val_loss: 0.0702 - val_accuracy: 0.9793
Epoch 2/5
750/750 [=====] - 11s 14ms/step - loss: 0.0556 - accuracy: 0.9826 - val_loss: 0.0499 - val_accuracy: 0.9856
Epoch 3/5
750/750 [=====] - 12s 16ms/step - loss: 0.0371 - accuracy: 0.9881 - val_loss: 0.0397 - val_accuracy: 0.9874
Epoch 4/5
750/750 [=====] - 11s 14ms/step - loss: 0.0289 - accuracy: 0.9906 - val_loss: 0.0368 - val_accuracy: 0.9886
Epoch 5/5
750/750 [=====] - 11s 14ms/step - loss: 0.0227 - accuracy: 0.9927 - val_loss: 0.0366 - val_accuracy: 0.9900
313/313 [=====] - 1s 4ms/step - loss: 0.0300 - accuracy: 0.9901
Test accuracy: 0.9901000261306763
```

Comparing CNN Implementations

This write-up compares two Convolutional Neural Network (CNN) implementations for a specific task. One approach builds the CNN from scratch, while the other leverages deep learning libraries.

Implementation Approaches:

From-Scratch CNN: This method involves manually coding all the mathematical operations and functionalities of a CNN architecture. It offers complete control but requires significant development effort and computational resources.

Deep Learning Library CNN: This method utilizes pre-built functions and modules offered by deep learning libraries like TensorFlow or PyTorch. We used Tensorflow here. This approach simplifies development, reduces coding time, and often provides optimized implementations.

Key Observations:

Accuracy: The deep learning library implementation achieved a higher accuracy (99.01%) compared to the from-scratch approach (97.98%). This could be due to several factors:

Speed: The write-up highlights that the deep learning library implementation was significantly faster. This is likely due to:

Pre-built Functions: Libraries offer optimized functions for convolutions, pooling, and other core operations, leading to faster execution. **Hardware Acceleration:** Libraries can leverage hardware acceleration capabilities of GPUs or TPUs for faster training and inference.

Trade-offs and Considerations:

Development Time: From-scratch implementations require more time and expertise. Libraries offer faster development cycles. **Control and Flexibility:** From-scratch approaches provide complete control over the architecture, but libraries offer flexibility through pre-trained models and modular components. **Computational Resources:** Both approaches require significant computational resources for training. Libraries might leverage hardware acceleration for efficiency. **Conclusion:**

This comparison showcases the potential benefits of using deep learning libraries for CNN development. Libraries offer faster development, potentially higher accuracy, and improved efficiency. However, from-scratch implementations can be valuable for gaining a deeper understanding of CNNs or for building custom architectures not available in libraries.

Choosing the Right Approach:

The choice between these approaches depends on specific project requirements. If development speed and high accuracy are priorities, deep learning libraries are a compelling option. However, if complete control over the architecture or a deep understanding of CNNs is crucial, a from-scratch approach might be considered.

Q3. Apply five different regularisation methods and comment on the performance of deep learning data. You can use MNIST data or any other of your choice. You may use the architecture designed in Q2 or may take the existing methods like VGG/ ResNet . (10)

```
In [1]: import tensorflow as tf
from tensorflow.keras.datasets import mnist
from matplotlib import pyplot as plt
(trainX, trainy), (testX, testy) = mnist.load_data()
```

Regularizer 1: L1 regulariser

```
In [11]: # Define the CNN architecture
from keras.regularizers import l1

model = models.Sequential([
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1), kernel_regularizer=l1),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu', kernel_regularizer=l1),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu', kernel_regularizer=l1),
    layers.Flatten(),
    layers.Dense(64, activation='relu'),
    layers.Dense(10, activation='softmax')
])

# Compile the model
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

# Train the model
history = model.fit(train_images, train_labels, epochs=5, batch_size=64, validation_data=(test_images, test_labels))

# Evaluate the model
test_loss, test_acc = model.evaluate(test_images, test_labels)
print(f'Test accuracy: {test_acc}')
```

Epoch 1/5
750/750 [=====] - 11s 14ms/step - loss: 0.2340 - accuracy: 0.9359 - val_loss: 0.0945 - val_accuracy: 0.9768
Epoch 2/5
750/750 [=====] - 10s 14ms/step - loss: 0.0796 - accuracy: 0.9803 - val_loss: 0.0785 - val_accuracy: 0.9815
Epoch 3/5
750/750 [=====] - 11s 14ms/step - loss: 0.0575 - accuracy: 0.9871 - val_loss: 0.0611 - val_accuracy: 0.9868
Epoch 4/5
750/750 [=====] - 11s 14ms/step - loss: 0.0475 - accuracy: 0.9893 - val_loss: 0.0627 - val_accuracy: 0.9850
Epoch 5/5
750/750 [=====] - 11s 15ms/step - loss: 0.0395 - accuracy: 0.9907 - val_loss: 0.0500 - val_accuracy: 0.9890
313/313 [=====] - 1s 4ms/step - loss: 0.0399 - accuracy: 0.9913
Test accuracy: 0.9912999868392944

Regularizer 2: L2 regulariser


```
In [12]: # Define the CNN architecture
from keras.regularizers import l2

model = models.Sequential([
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1), kernel_regularizer=l2(0.01)),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu', kernel_regularizer=l2(0.01)),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu', kernel_regularizer=l2(0.01)),
    layers.Flatten(),
    layers.Dense(64, activation='relu'),
    layers.Dense(10, activation='softmax')
])

# Compile the model
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

# Train the model
history = model.fit(train_images, train_labels, epochs=5, batch_size=64, validation_data=(test_images, test_labels))

# Evaluate the model
test_loss, test_acc = model.evaluate(test_images, test_labels)
print(f'Test accuracy: {test_acc}')
```

```
Epoch 1/5
750/750 [=====] - 12s 15ms/step - loss: 0.2147 - accuracy: 0.9370 - val_loss: 0.0850 - val_accuracy: 0.9753
Epoch 2/5
750/750 [=====] - 11s 15ms/step - loss: 0.0615 - accuracy: 0.9815 - val_loss: 0.0567 - val_accuracy: 0.9843
Epoch 3/5
750/750 [=====] - 11s 14ms/step - loss: 0.0455 - accuracy: 0.9870 - val_loss: 0.0439 - val_accuracy: 0.9883
Epoch 4/5
750/750 [=====] - 11s 15ms/step - loss: 0.0342 - accuracy: 0.9902 - val_loss: 0.0500 - val_accuracy: 0.9865
Epoch 5/5
750/750 [=====] - 10s 13ms/step - loss: 0.0304 - accuracy: 0.9912 - val_loss: 0.0469 - val_accuracy: 0.9867
313/313 [=====] - 1s 3ms/step - loss: 0.0408 - accuracy: 0.9883
Test accuracy: 0.9883000254631042
```

Regularizer 3: Dropout

```

In [5]: # Define the CNN architecture
from keras.regularizers import l2

model = models.Sequential([
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.Dropout(0.25), # Add dropout after last convolutional layer
    layers.Flatten(),
    layers.Dense(64, activation='relu'),
    layers.Dropout(0.25), # Add dropout after first dense layer
    layers.Dense(10, activation='softmax')
])

# Compile the model
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

# Train the model
history = model.fit(train_images, train_labels, epochs=5, batch_size=64, val

# Evaluate the model
test_loss, test_acc = model.evaluate(test_images, test_labels)
print(f'Test accuracy: {test_acc}')

```

```

Epoch 1/5
750/750 [=====] - 12s 15ms/step - loss: 0.2855 - a
ccuracy: 0.9107 - val_loss: 0.0587 - val_accuracy: 0.9818
Epoch 2/5
750/750 [=====] - 11s 15ms/step - loss: 0.0818 - a
ccuracy: 0.9752 - val_loss: 0.0526 - val_accuracy: 0.9831
Epoch 3/5
750/750 [=====] - 11s 15ms/step - loss: 0.0604 - a
ccuracy: 0.9818 - val_loss: 0.0420 - val_accuracy: 0.9871
Epoch 4/5
750/750 [=====] - 12s 16ms/step - loss: 0.0522 - a
ccuracy: 0.9840 - val_loss: 0.0528 - val_accuracy: 0.9851
Epoch 5/5
750/750 [=====] - 11s 15ms/step - loss: 0.0423 - a
ccuracy: 0.9875 - val_loss: 0.0375 - val_accuracy: 0.9898
313/313 [=====] - 1s 3ms/step - loss: 0.0293 - acc
uracy: 0.9898
Test accuracy: 0.989799976348877

```

Regularizer 4: Batch normalisation

```
In [6]: from keras.layers import BatchNormalization

# Define the CNN architecture

model = models.Sequential([
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)),
    layers.BatchNormalization(), # Add BatchNormalization after Conv2D
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.BatchNormalization(), # Add BatchNormalization after Conv2D
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.BatchNormalization(), # Add BatchNormalization after Conv2D
    layers.Flatten(),
    layers.Dense(64, activation='relu'),
    layers.BatchNormalization(), # Add BatchNormalization after Conv2D
    layers.Dense(10, activation='softmax')
])

# Compile the model
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

# Train the model
history = model.fit(train_images, train_labels, epochs=5, batch_size=64, val

# Evaluate the model
test_loss, test_acc = model.evaluate(test_images, test_labels)
print(f'Test accuracy: {test_acc}')
```

```
Epoch 1/5
750/750 [=====] - 15s 19ms/step - loss: 0.1236 - a
ccuracy: 0.9640 - val_loss: 0.0537 - val_accuracy: 0.9845
Epoch 2/5
750/750 [=====] - 14s 18ms/step - loss: 0.0418 - a
ccuracy: 0.9874 - val_loss: 0.0455 - val_accuracy: 0.9863
Epoch 3/5
750/750 [=====] - 14s 19ms/step - loss: 0.0305 - a
ccuracy: 0.9902 - val_loss: 0.0462 - val_accuracy: 0.9857
Epoch 4/5
750/750 [=====] - 13s 18ms/step - loss: 0.0213 - a
ccuracy: 0.9934 - val_loss: 0.0434 - val_accuracy: 0.9873
Epoch 5/5
750/750 [=====] - 15s 20ms/step - loss: 0.0173 - a
ccuracy: 0.9946 - val_loss: 0.0653 - val_accuracy: 0.9792
313/313 [=====] - 1s 4ms/step - loss: 0.0764 - acc
uracy: 0.9753
Test accuracy: 0.9753000140190125
```

Regularizer 5: Early stop

```
In [7]: from tensorflow.keras.callbacks import EarlyStopping
# Define the CNN architecture

model = models.Sequential([
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.Flatten(),
    layers.Dense(64, activation='relu'),
    layers.Dense(10, activation='softmax')
])

# Compile the model
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])
early_stopping = EarlyStopping(monitor='val_loss', patience=5) # Monitor va
# Train the model
history = model.fit(train_images, train_labels, epochs=5, batch_size=64, val

# Evaluate the model
test_loss, test_acc = model.evaluate(test_images, test_labels)
print(f'Test accuracy: {test_acc}')
```

```
Epoch 1/5
750/750 [=====] - 13s 17ms/step - loss: 0.2280 - a
ccuracy: 0.9309 - val_loss: 0.0757 - val_accuracy: 0.9787
Epoch 2/5
750/750 [=====] - 12s 16ms/step - loss: 0.0593 - a
ccuracy: 0.9822 - val_loss: 0.0525 - val_accuracy: 0.9852
Epoch 3/5
750/750 [=====] - 11s 14ms/step - loss: 0.0421 - a
ccuracy: 0.9863 - val_loss: 0.0510 - val_accuracy: 0.9859
Epoch 4/5
750/750 [=====] - 10s 13ms/step - loss: 0.0330 - a
ccuracy: 0.9898 - val_loss: 0.0435 - val_accuracy: 0.9874
Epoch 5/5
750/750 [=====] - 10s 13ms/step - loss: 0.0260 - a
ccuracy: 0.9913 - val_loss: 0.0461 - val_accuracy: 0.9876
313/313 [=====] - 1s 3ms/step - loss: 0.0327 - acc
uracy: 0.9902
Test accuracy: 0.9901999831199646
```

Observed Performance

Overall Performance:

All regularization techniques achieved very high test accuracy (above 97.5%), indicating effective prevention of overfitting and good generalization on unseen data.

With L1 regularisation, we obtained an accuracy of 99.13 %, L2 regularisation gave 98.83 %. With dropout, batch normalisation and early stop, the accuracies were 98.98 %, 97.53 % and 99.02 % respectively.

Comparative Analysis:

L1 Regularization (highest: 99.13%) and Early Stopping (99.02%) emerged as the top performers in terms of accuracy. L1 might have led to a slightly sparser model with fewer influential weights, potentially reducing overfitting. Early stopping might have prevented the model from overtraining on the training data. Dropout (98.98%) performed very close to the leaders, suggesting its effectiveness in randomly dropping activations and preventing overfitting. L2 Regularization (98.83%) achieved slightly lower accuracy compared to L1, but it's still a strong contender. L2 tends to shrink weights towards zero, potentially leading to a smoother decision boundary but might be slightly less effective than L1 for achieving sparsity in some cases. Batch Normalization (97.53%) showed the lowest accuracy among the tested techniques. While it can improve training speed and stability, it might not have provided the strongest regularization effect in this specific scenario.

Important Considerations:

It's important to note that these results are based on a single test run. Repeating the experiment with different random seeds could lead to slight variations in accuracy. The optimal choice of regularization technique can depend on factors like the dataset size, complexity, and noise levels. You might also consider other metrics beyond accuracy, such as validation loss or training time, when making a final decision.

Recommendations:

Given the close performance between L1 and Early Stopping, you can experiment further to see which one is more consistent across different random seeds. L2 and Dropout are still strong options, and the choice might depend on whether you prefer weight shrinkage (L2) or promoting sparsity (L1). Consider running the model for a longer duration (more epochs) with Early Stopping to see if it can reach even higher accuracy. Batch Normalization might be more beneficial for deeper networks or for datasets with internal covariate shift. Remember, the best approach depends on your specific dataset and goals. It's always recommended to experiment with different techniques and hyperparameter settings to find the optimal configuration for your CNN model.